

Final Assignment - Bayesian Statistics

Nina van Gerwen (1860852)

15/06/2022

Introduction

Bayesian Statistics is a field of statistics that is based on another interpretation of probabilities. In Bayesian Statistics, a probability can be viewed as a quantification of personal belief in an event (source). In the current paper, we will investigate the predictive ability of Social Economical Status (SES) and Verbal Ability (VA) on intelligence as measured by the Intelligence Quotient (IQ) through Bayesian methods and statistics.

Research Questions

In this paper, we will aim to research whether IQ can be predicted by both SES and VA. To answer this research question, we have come up with a total of three hypotheses.

Hypotheses: - IQ in students is predicted by both their Verbal Ability, while statistically controlling for SES, and their SES, while controlling for their Verbal Ability, where people with a higher Verbal Ability or SES will have higher IQ scores.

$$H_1 : \beta_1 > 0, \beta_2 > 0$$

- IQ in students is predicted only by their Verbal Ability, while statistically controlling for SES, where people with a higher Verbal Ability will have higher IQ scores.

$$H_2 : \beta_1 > 0, \beta_2 \approx 0$$

- IQ in students is predicted only by their SES, while statistically controlling for Verbal Ability, where people with a higher self-reported SES will have higher predicted IQ scores.

$$H_3 : \beta_1 \approx 0, \beta_2 > 0$$

OPTIONAL INTERACTION EFFECT - The effect of Verbal Ability on IQ depends on self-reported SES, where people who have a higher self-reported SES will have a greater effect of Verbal Ability on IQ than people with lower self-reported SES scores.

$$H_3 : \beta_1, \beta_2, \beta_3 > 0$$

OPTIONAL: one $>$, other unconstrained

Method

Description of the data

For our dataset, we chose one gained from Multivariate Statistics. The dataset contains a total of three variables from 400 high school students and contains no missing data. The variables are Verbal Ability, which ranges from 2.26 to 4 ($M = 3.39$, $SD = 0.38$), self-reported SES on a scale from 0 to 60 ($M = 32.78$, $SD = 9.66$) and IQ, which ranged from 85.5 to 143 ($M = 118.2$, $SD = 11.92$). When visually inspecting the histograms of the variables, we find that SES and IQ seem to follow a normal distribution. Verbal Ability, however, seems to be slightly skewed to the left which could be explained by a ceiling effect. For all analyses, we grand mean centered the variable of Verbal Ability in order to aid convergence and interpretation.

Statistical Analyses

To answer our research questions and hypotheses, we will run a total of three Bayesian linear regression analyses. In the first two analyses, we will use Gibbs Sampling with both an uninformative (LR1) and informative (LR2) prior for the estimation of our regression coefficients. For the final analysis, regression coefficients will be estimated through an independent Metropolis-Hastings (MH) Sampler (LR3), where the proposal distribution is a Student's t -distribution with 1 degree of freedom in order to maximize the uncertainty in the tails.

Prior distributions

IQ (b1): known to be a variable that has a mean of 100, sd of 15. SES (b2): check categorical distribution in society and try to mimic that in the self-report way variance (vari): intercept??

- nvm, need priors for mu/sigma.. not their distribution (LEL)

Model Diagnostics

First, we will assess the convergence of all analyses through visual inspection of both traceplots and autocorrelation plots. (Because the analyses consist of only one chain, we can not use the Gelman-Rubin statistic. Furthermore, the MC error is based on a test and visual tests are preferred)

Second, we will assess whether the different methods of sampling gave different results. If they do not seem to differ, we will try to choose the best model to continue with.

– other things to say perhaps only continue with the results gained from Gibbs Samplers with informative priors. (reasons)

Finally, for the chosen model, we will assess whether the residuals of our model are normally distributed through use of the posterior predictive p -value and discrepancy measures with the following original test statistic:

$$T = (\mu - m_1)^2$$

where μ is the mean and m_1 is the median of the distribution of the residuals. The reasoning behind the statistic is that the more skewed a distribution is, the larger the difference between the mean and median will be. However, this difference can be either negative (in a left-skewed distribution) or positive (in a right-skewed distribution). Hence, we square the difference in order to always gain a positive value. Now, the larger the statistic is, the larger the absolute difference between mean and median, and the more skewed a distribution supposedly is.

To investigate whether our original statistic works, we will compare the results to a posterior predictive p -value with the known statistic of skewness:

$$\tilde{\mu}_3 = \frac{\sum_i^N (X_i - \bar{X})^3}{(N - 1) \cdot \sigma^3}$$

where \bar{X} is the mean of the residuals and σ the standard deviation of the residuals. Skewness is known to measure the asymmetry of a distribution. Therefore, we will use it to see whether the results are similar to the results of our original test statistic.

Model comparison

Besides answering the hypotheses, we also want to see which hypothesis has the most support. To do this, we will compare the hypotheses (and their corresponding statistical models) through the use of the Bayes factor. When comparing the hypotheses, their priors distributions will be fractional, where the size of the fraction is decided by the following formula:

$$b = \frac{J}{N}$$

where b is the denominator of the fraction, J is the number of independent constraints in the hypothesis and N is sample size. However, we will conduct a sensitivity analysis with varying numbers of J to see the influence it has on the Bayes Factor.

Results

Parameter estimates

For our LR1 analysis, we found an intercept of 118.32 (95% *C.I.*: [114.49; 122.14]). For the effect of Verbal Ability on IQ, we found a regression coefficient of 11.93 (95% *C.I.*: [9.08; 14.77]) and for the effect of SES on IQ, we found a regression coefficient of 0.003 (95% *C.I.*: [-0.11; 0.12]).

– interpretation –

For our LR2 analysis, we found XXX.

– interpretation –

For the final LR3 analysis, we found very similar results to the first analysis with an intercept of 115.9261 (95% *C.I.*: [114.47; 122.21]) and two regression coefficients of 11.91 (95% *C.I.*: [9.11; 14.77]) and 0.083 (95% *C.I.*: [-0.11; 0.11]).

– interpretation –

Model diagnostics

When investigating the convergence of the three analyses, we found the following results. For LR1, the traceplots of all coefficients showed very sound convergence. The autocorrelation plots showed two notable results. For the intercept and regression coefficient for SES, the autocorrelation started at high values around .90 and then in a monotone decreasing fashion reached 0 at lag 45. The other autocorrelation plots remained constant at values around 0 for every lag. Although the autocorrelation could be improved upon, we feel it is still safe to assume the first model converged properly when also taking the traceplot into account.

The second analysis, LR2, ...

For the final analysis, LR3, the traceplots for the coefficients of the independent variables showed sound convergence. The traceplot for the intercept, while still acceptable, showed that sometimes quite extreme values were accepted and convergence of this coefficient could be improved. The traceplot of the variance, however, showed that there might have been some issues as the plot shows that somehow it wanted to accept negative values. Furthermore, some very extreme values were accepted, which might lead to issues if we were to use this model for model assumption checks through the posterior predictive p -value. The autocorrelation plots all showed constant values close to 0. Hence, taken together, we assume that the model converged adequately. However, due to the extreme values, we refrained from using this model to check the normality of the residuals through use of the posterior predictive p -value.

Checking the assumption of normality of residuals, we found a posterior predictive p -value of $<.001$ for both our original and skewness statistic. These p -values mean the following: for our original statistic, it entails that when using the same sampled coefficients from their posterior distribution to calculate the residuals, we find a larger difference between the mean and median of these residuals when using our observed dataset compared to using a simulated dataset. For the skewness statistic, it means also means that the calculated residuals in our observed dataset were practically always more skewed than the residuals in a simulated dataset.

From these results, we can conclude that the residuals most likely do not follow a normal distribution and instead are skewed. This result is further strengthened by visual inspection of the residuals, which shows that the distribution is slightly skewed to the right. skew is unknown.

Model comparison

Discussion

- Differences between Bayesian and Frequentist analyses (relevant to the context of the research questions)